Information Retrieval

Homework 2

1. **Introduction**

   The objective of this assignment is to calculate the term weights for the tokens that occur in each document in the collection. To execute the program, it should be sufficient to type:

   python3 calcwts.py input-directory output-directory

   where, the input-directory contains the original files (i.e., the unprocessed files directory) and the output-directory contains all generated output files.

   The program also needs a temporary folder with the name "temp" to store intermediate results after preprocessing

2. **Input**

   The input directory contains a set of HTML documents that have to be processed

3. **Output**

   The goal is to build one output file per input file. In the output file, there would be one line per token that survives preprocessing. That line would contain the token and the token weight.

4. **Preprocessing**

   I have used the processor from Homework1 to process the documents. After loading the file contents, BeautifulSoup library is used to parse the HTML documents and extract information in pure text format. Any special characters are removed using the concept of regular expressions. NLTK library is used to tokenize the file content which results in a list containing individual words from each particular document.

   The initial processor is further extended to do the following:

   a) Remove words of length 1

   b) Remove all the words occurring in the file stopwords.txt

   c) Remove words occurring only once in the entire corpus

   Figure below shows two HTML files before and after preprocessing.

Fig 1 : Left box has part of HTMl1 before preprocessing

Right box has part of tokenized words and their frequencies

Fig 2: Left box has part of HTML2 before preprocessing

Right box has part of tokenized words and their frequencies

5. Term Weighting

Okapi BM25 function is used to calculate term weights.It is a ranking function used by search engines to rank matching documents according to their relevance to a given search query. Given a query Q, containing keywords q_1, ..., q_n, the BM25 score of

a document D is

$$\text{score}(D, Q) = \sum_{i=1}^{n} \text{IDF}(q_i) \cdot \frac{f(q_i, D) \cdot (k_1 + 1)}{f(q_i, D) + k_1 \cdot \left(1 - b + b \cdot \frac{|D|}{\text{avgdl}}\right)},$$

where f(q{i},D)) is term frequency in the document D, |D| is the length of the document D in words, and avgdl is the average document length in the text collection from which documents are drawn.

K1 and b are free parameters whose values have been chosen as 1.2 and 0.75 respectively.

Inverse document frequency(idf) is calculated using the formula log(N/n(q)),

Where N is the total number of documents in the corpus and n(q) is the number of documents containing the word q.Since the above BM25 formula uses length of document D as well as the average length of documents, it also takes the normalization factor into account. In the final output files, the term weights of some words are shown to be negative. This is because the BM25 function calculates term weights of words that occur in more than half of the document collection to be negative.

publishers 3.22	deputy 3.33
employees 2.1	posted 4.61
california 3.56	obtained 4.88
1977 4.61	share 4.61
edition 4.55	prescription 5.53
bribe 3.78	employee 2.24
star 4.61	goals 3.86
newspaper 0.6	realize 4.61
role 1.43	organizational 5.53
police 2.1	efforts 2.77
viewpoints 4.61	findings 4.61
established 3.79	insurers 4.88
threatened 2.77	extent 2.41
change 2.2	penalties 4.88
national 1.6	shield 4.61
gov -0.81	seen 3.22
pressure 1.39	demands 5.53
eventually 4.61	interactive 3.79
despite 2.77	remain 3.79
decided 4.55	participation 5.53
journalists -2.14	despite 2.77
war 3.79	course 3.79
columnist 5.53	ultimately 3.79
instance 1.83	differing 4.61
bound 4.61	hold 3.22
shooting 2.77	governing 5.53

Fig 3 : Term weights of a set of words from documents HTML1 and HTML2

6. Evaluation

The program has been tested on different number of files and the graph shown below has been plotted by indexing time as a function of the number of documents. The timings shown is for the entire process of reading from files , preprocessing , calculating term weights and writing to output files.

Tokenization and calculating term weights are the two functions that utilized most of the time.

Also, since I have used a temporary folder to store intermediate results from tokenization, this has added an extra step of reading and writing to files which contributes to the significant increase in time as the number of files increases.


Fig 4: Execution time vs number of files